

# Machine Learning Online Class

Exercise 1: Linear regression with multiple variables

Instructions

-----

This file contains code that helps you get started on the linear regression exercise.

You will need to complete the following functions in this exercise:

```
warmUpExercise.m
plotData.m
gradientDescent.m
computeCost.m
gradientDescentMulti.m
computeCostMulti.m
featureNormalize.m
normalEqn.m
```

For this part of the exercise, you will need to change some parts of the code below for various experiments (e.g., changing learning rates).

## Initialization

### ===== Part 1: Feature Normalization =====

## Clear and Close Figures

```
clear ; close all; clc

fprintf('Loading data ...\n');
```

Loading data ...

## Load Data

```
data = load('ex1data2.txt');
X = data(:, 1:2);
y = data(:, 3);
m = length(y);

% Print out some data points
fprintf('First 10 examples from the dataset: \n');
```

First 10 examples from the dataset:

```
fprintf(' x = [%.0f %.0f], y = %.0f \n', [X(1:10,:) y(1:10,:)]);
```

```
x = [2104 3], y = 399900
x = [1600 3], y = 329900
x = [2400 3], y = 369000
x = [1416 2], y = 232000
x = [3000 4], y = 539900
```

```

x = [1985 4], y = 299900
x = [1534 3], y = 314900
x = [1427 3], y = 198999
x = [1380 3], y = 212000
x = [1494 3], y = 242500

```

```
fprintf('Program paused. Press enter to continue.\n');
```

Program paused. Press enter to continue.

```
pause;
```

```

% Scale features and set them to zero mean
fprintf('Normalizing Features ...\n');

```

Normalizing Features ...

```
[X mu sigma] = featureNormalize(X);
```

```

% Add intercept term to X
X = [ones(m, 1) X];

```

## ===== Part 2: Gradient Descent =====

```

% ===== YOUR CODE HERE =====
% Instructions: We have provided you with the following starter
%               code that runs gradient descent with a particular
%               learning rate (alpha).
%
%               Your task is to first make sure that your functions -
%               computeCost and gradientDescent already work with
%               this starter code and support multiple variables.
%
%               After that, try running gradient descent with
%               different values of alpha and see which one gives
%               you the best result.
%
%               Finally, you should complete the code at the end
%               to predict the price of a 1650 sq-ft, 3 br house.
%
% Hint: By using the 'hold on' command, you can plot multiple
%       graphs on the same figure.
%
% Hint: At prediction, make sure you do the same feature normalization.
%

```

```
fprintf('Running gradient descent ...\n');
```

Running gradient descent ...

```

% Choose some alpha value
alpha = 0.01;

```

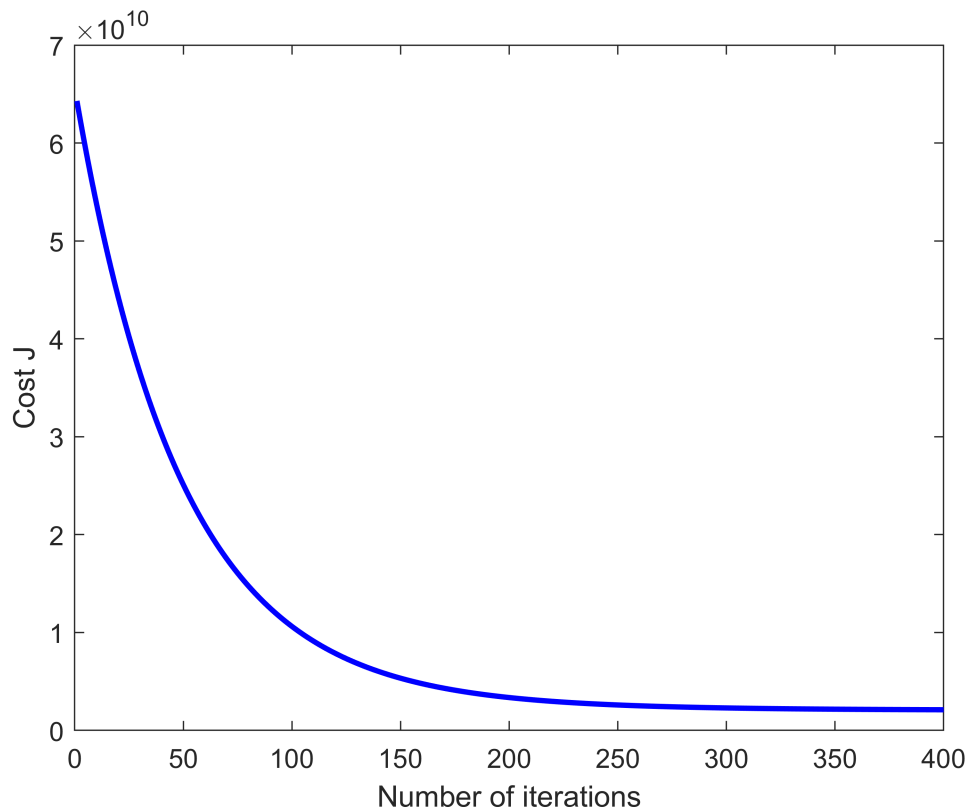
```

num_iters = 400;

% Init Theta and Run Gradient Descent
theta = zeros(3, 1);
[theta, J_history] = gradientDescentMulti(X, y, theta, alpha, num_iters);

% Plot the convergence graph
figure;
plot(1:numel(J_history), J_history, '-b', 'LineWidth', 2);
xlabel('Number of iterations');
ylabel('Cost J');

```



```

% Display gradient descent's result
fprintf('Theta computed from gradient descent: \n');

```

Theta computed from gradient descent:

```
fprintf(' %f \n', theta);
```

```

334302.063993
100087.116006
3673.548451

```

```
fprintf('\n');
```

```

% Estimate the price of a 1650 sq-ft, 3 br house
% ===== YOUR CODE HERE =====
% Recall that the first column of X is all-ones. Thus, it does
% not need to be normalized.

```

```
price = [1, (([1650, 3] - mu) ./ sigma)] * theta; % You should change this
```

```
% =====
```

```
fprintf(['Predicted price of a 1650 sq-ft, 3 br house ' ...  
        '(using gradient descent):\n $%f\n'], price);
```

```
Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):  
$289314.620338
```

```
fprintf('Program paused. Press enter to continue.\n');
```

```
Program paused. Press enter to continue.
```

```
pause;
```

## =====**Part 3: Normal Equations**=====

```
fprintf('Solving with normal equations...\n');
```

```
Solving with normal equations...
```

```
% ===== YOUR CODE HERE =====  
% Instructions: The following code computes the closed form  
%               solution for linear regression using the normal  
%               equations. You should complete the code in  
%               normalEqn.m  
%  
%               After doing so, you should complete this code  
%               to predict the price of a 1650 sq-ft, 3 br house.  
%  
%
```

## Load Data

```
data = csvread('ex1data2.txt');  
X = data(:, 1:2);  
y = data(:, 3);  
m = length(y);
```

```
% Add intercept term to X  
X = [ones(m, 1) X];
```

```
% Calculate the parameters from the normal equation  
theta = normalEqn(X, y);
```

```
% Display normal equation's result  
fprintf('Theta computed from the normal equations: \n');
```

```
Theta computed from the normal equations:
```

```
fprintf(' %f \n', theta);
```

```
89597.909544
139.210674
-8738.019113
```

```
fprintf('\n');

% Estimate the price of a 1650 sq-ft, 3 br house
% ===== YOUR CODE HERE =====
price = [1,1650,3] * theta; % You should change this

% =====

fprintf(['Predicted price of a 1650 sq-ft, 3 br house ' ...
        '(using normal equations):\n $%f\n'], price);
```

```
Predicted price of a 1650 sq-ft, 3 br house (using normal equations):
$293081.464335
```